

In []:

```
In [13]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, accuracy_score
import statsmodels.api as sm
```

```
In [14]: std_prfmnce = pd.read_csv("C:/Users/ardra/Downloads/archive (14)/StudentPerforman
std_prfmnce
```

Out[14]:

	Hours_Studied	Attendance	Parental_Involvement	Access_to_Resources	Extracurri
0	23	84	Low	High	
1	19	64	Low	Medium	
2	24	98	Medium	Medium	
3	29	89	Low	Medium	
4	19	92	Medium	Medium	
...
6602	25	69	High	Medium	
6603	23	76	High	Medium	
6604	20	90	Medium	Low	
6605	10	86	High	High	
6606	15	67	Medium	Low	

6607 rows × 20 columns

In [15]: `print(std_prfmnce.columns)`

```
Index(['Hours_Studied', 'Attendance', 'Parental_Involvement',
      'Access_to_Resources', 'Extracurricular_Activities', 'Sleep_Hours',
      'Previous_Scores', 'Motivation_Level', 'Internet_Access',
      'Tutoring_Sessions', 'Family_Income', 'Teacher_Quality', 'School_Type',
      'Peer_Influence', 'Physical_Activity', 'Learning_Disabilities',
      'Parental_Education_Level', 'Distance_from_Home', 'Gender',
      'Exam_Score'],
      dtype='object')
```

```
In [16]: print("\nDescriptive Statistics:")
print(std_prfmnce.describe())
```

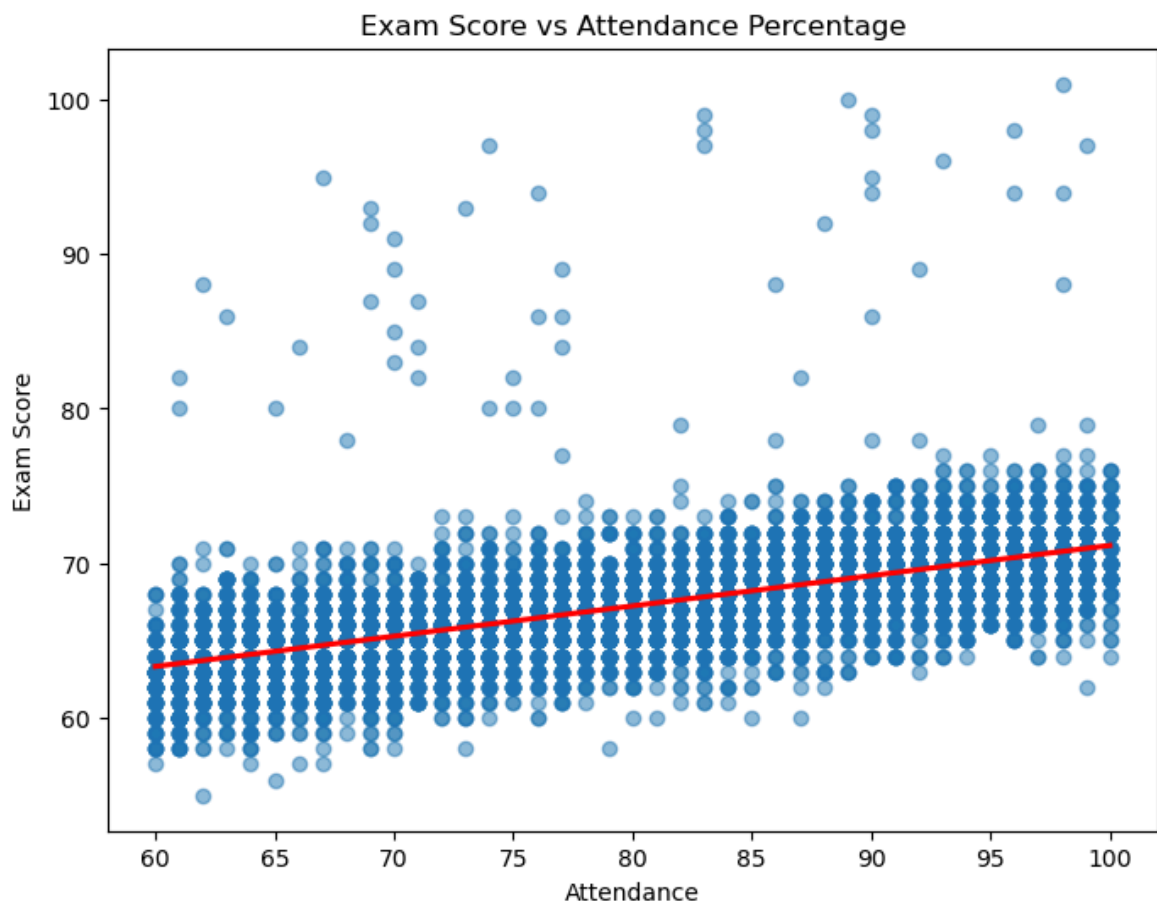
Descriptive Statistics:

	Hours_Studied	Attendance	Sleep_Hours	Previous_Scores \
count	6607.000000	6607.000000	6607.000000	6607.000000
mean	19.975329	79.977448	7.02906	75.070531
std	5.990594	11.547475	1.46812	14.399784
min	1.000000	60.000000	4.00000	50.000000
25%	16.000000	70.000000	6.00000	63.000000
50%	20.000000	80.000000	7.00000	75.000000
75%	24.000000	90.000000	8.00000	88.000000
max	44.000000	100.000000	10.00000	100.000000

	Tutoring_Sessions	Physical_Activity	Exam_Score
count	6607.000000	6607.000000	6607.000000
mean	1.493719	2.967610	67.235659
std	1.230570	1.031231	3.890456
min	0.000000	0.000000	55.000000
25%	1.000000	2.000000	65.000000
50%	1.000000	3.000000	67.000000
75%	2.000000	4.000000	69.000000
max	8.000000	6.000000	101.000000

```
In [17]: plt.figure(figsize=(8, 6))
sns.regplot(x=std_prfmnce['Attendance'], y=std_prfmnce['Exam_Score'], scatter_kw=
plt.xlabel("Attendance")
plt.ylabel("Exam Score")
plt.title("Exam Score vs Attendance Percentage")
plt.show()

print("The correlation between the two variables is")
print(std_prfmnce['Attendance'].corr(std_prfmnce['Exam_Score']))
```



The correlation between the two variables is
0.5810718633120644

```
In [20]: # Prepare Data for Regression
X = std_prfmnce[['Attendance']] # Independent variable
Y = std_prfmnce['Exam_Score'] # Dependent variable

X = sm.add_constant(X)
X.head()
model = sm.OLS(Y,X, missing='drop')
model_result=model.fit()
model_result.summary()
```

Out[20]:

OLS Regression Results


Dep. Variable:	Exam_Score	R-squared:	0.338			
Model:	OLS	Adj. R-squared:	0.338			
Method:	Least Squares	F-statistic:	3367.			
Date:	Sat, 15 Mar 2025	Prob (F-statistic):	0.00			
Time:	20:42:32	Log-Likelihood:	-16989.			
No. Observations:	6607	AIC:	3.398e+04			
Df Residuals:	6605	BIC:	3.400e+04			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	51.5786	0.273	189.191	0.000	51.044	52.113
Attendance	0.1958	0.003	58.026	0.000	0.189	0.202
Omnibus:	4752.483	Durbin-Watson:	2.023			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	163501.063			
Skew:	3.038	Prob(JB):	0.00			
Kurtosis:	26.601	Cond. No.	566.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [22]: fitted_values = model_result.fittedvalues
residuals = model_result.resid

plt.figure(figsize=(8,5))
plt.scatter(fitted_values, residuals, alpha=0.8)
plt.axhline(y=0, color='r', linestyle='dashed')
plt.xlabel('fitted_values')
plt.ylabel('Residuals')
plt.title('Residuals Vs. Fitted Values')
plt.show()
```



```
In [24]: import statsmodels.stats.diagnostic as smd
import numpy as np

# Get residuals and independent variables (exog)
residuals = model_result.resid
exog = model_result.model.exog # Independent variables (including constant)

# Perform the Breusch-Pagan test
bp_test = smd.het_breuschpagan(residuals, exog)
p_value = bp_test[1] # Extract the p-value

print(f"Breusch-Pagan test p-value: {p_value}")

# Interpretation
if p_value < 0.05:
    print("Heteroscedasticity detected! Consider transformation or robust standa
else:
    print("No significant heteroscedasticity detected.")
```

Breusch-Pagan test p-value: 0.8057839143058365

No significant heteroscedasticity detected.

In []: